WORKSHOP

Trustworthy AI for Good @ NeurIPS 2026

Date: TBD | Location: Paris, France

Submit on OpenReview   I'd Like to Be a Reviewer: Register Here!   View ICML 2026 Workshop

Sponsors: We're open for sponsorship with PR benefits! Please email zjingchen@cs.toronto.edu directly!

Workshop Overview

Agentic AI systems increasingly shape how billions of people engage with public institutions, civic discourse, and society at large. While much work has focused on making models safer in avoiding harmful output, it is equally important for these improvements to translate into social good at scale.

The **AI4GOOD** workshop brings together the **AI safety**, **AI for social good**, and **AI policy/governance** communities to connect what models can do as individual systems with what they do when deployed across populations. We aim to bridge technical advances in trustworthy AI with real-world societal impact, including protecting democratic institutions and civic discourse.

Call for Papers

We welcome submissions in a wide range of topics (if you're not sure about your paper, we encourage you to just submit!).

Trustworthy AI models: evaluation, auditing, and red-teaming of models for harmful behaviors and failure modes; safety monitoring after deployment; and alignment and robustness methods.

AI for social good: methods, evidence standards, and evaluation frameworks for demonstrating real-world benefit and avoiding unintended harms at population scale.

AI for information integrity: detection and mitigation of disinformation, manipulation, and influence operations; and building resilience of the information ecosystem.

AI for public institutions: accountable use of AI in government and civic settings, including transparency, documentation, procurement, and oversight practices.

AI for civic discourse: AI systems that support public deliberation and civic engagement while preserving legitimacy and avoiding undue influence.

Cooperative AI: multi-agent coordination, negotiation, and conflict resolution; and mechanisms for trust, commitment, and cooperation among AI systems and between AI and humans.

MULTI-AGENT TRACK

Multi-agent Security and Safety

This track will be co-hosted by Klaudia Krawiecka and Swapneel Mehta.

Topics include collusion and steganographic communication between agents, secure interaction and delegation protocols, confinement of strategic learned agents, compositional failures of safety and compliance guarantees, attacks that propagate across agent networks, evaluation and red-teaming of agent populations, and oversight and incident response for deployed multi-agent systems.

Read full call for papers

Submission Guidelines

- ✓ **Format:** Papers should be 2 to 9 pages (excluding references and appendices) using the NeurIPS 2026 workshop style.
- ✓ **Review Process:** All submissions will be reviewed double-blind via OpenReview. Please ensure your submission is fully anonymized.
- ✓ **Non-Archival:** Work may be submitted to or published at other venues.

Important Dates

KEY DATES

All deadlines are **11:59 PM Anywhere on Earth (AoE)** unless otherwise noted.

- 30 Jul 2026
  Submission Portal Opens
- Upcoming
  Paper Submission Deadline
  29 Aug 2026
- 29 Sep 2026
  Author Notification
- Camera Ready Deadline
  29 Nov 2026
- 12 Dec 2026
  Workshop Date

Full Call for Papers →

Speakers

Yoshua Bengio
University of Montreal & Mila & LawZero

Pavel Izmailov
New York University & Anthropic

Joelle Pineau
McGill University & Mila & Cohere

Been Kim
Google DeepMind

Sanmi Koyejo
Stanford University

John Sotiropoulos
Deep Cyber & OWASP

Workshop Schedule

The detailed schedule will be posted closer to the workshop date.

Organizing Team

Terry Jingchen Zhang
University of Oxford & Vector Institute

Changling Li
ETH Zürich & MPI-IS

Arian Khorasani
University of Toronto & Vector Institute

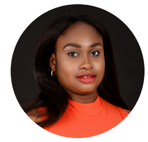

Joan Nwatu
University of Michigan

Yejin Son
University of British Columbia & Vector Institute

He (Shawn) Shuang
Palo Alto Networks

Jerick Shi
Carnegie Mellon University

Prakhar Gupta
University of Michigan

Kexin Li
University of Toronto & Schwartz Reisman Graduate Fellow

Wenjun Qiu
University of Toronto & Schwartz Reisman Graduate Fellow

Zhijing Jin
University of Toronto

Rada Mihalcea
University of Michigan

Milind Tambe
Harvard University

David Lie
University of Toronto & Schwartz Reisman Institute

Christian Schroeder de Witt
University of Oxford

Multi-Agent Track Lead

Klaudia Krawiecka
Meta

Swapneel Mehta
MIT

Frequently Asked Questions

I'm not sure if my paper is in scope

We encourage you to submit anyway! It is possible that we may have to desk-reject some papers that we believe might be a better fit for other venues. This will not be made public, so don't worry: it does not represent any judgment of your work, just an administrative necessity as our review capacity is limited.

Can I submit a paper that is under review at NeurIPS or another conference/journal?

On our side, we are open to any submissions currently under review (at NeurIPS or other venues). However, it is your responsibility to make sure that the other venue is OK with your submission to our workshop.

Can I submit to another NeurIPS workshop as well?

We are also open to that from our side, especially if you feel like your work spans the interest of several workshops. Similarly, you may want to check the website of other workshops to make sure they are OK with it as well.

Do I need to attend in person?

Details on presentation modality will be announced closer to the workshop date. We strongly encourage in-person presence when possible.