

Call for Papers

AI for Meta-Science

NEURIPS 2026 WORKSHOP · PARIS

[About](#)
[Call for Papers](#)
[Speakers](#)
[Schedule](#)

Overview and Scope

The AI for Meta-Science workshop concerns the use, evaluation, and governance of AI systems that shape the scientific process itself: peer review, reproducibility, research evaluation, publication incentives, and funding and governance mechanisms. We invite submissions to two tracks: a *technical track* for new methods, systems, and empirical analyses, and a *position track* for arguments about how the scientific ecosystem should change. The workshop is **non-archival**: submitting here leaves you free to publish the work at an archival venue later.

Tracks and Topics

Technical Track

New methods, systems, benchmarks, datasets, and empirical analyses of AI in the scientific process. Topics include, but are not limited to:

Position Track

Position papers argue for a viewpoint about what should be done, in contrast to technical papers, which report on completed advances. Submissions

- AI/LLM-assisted peer review: methods, evaluation, real-world deployments, and failure modes
- Benchmarks and evaluation of AI scientists and autonomous research agents
- Verification, reproducibility, and quality control of (AI-generated) scientific artifacts
- Detection of AI-generated content and robustness of review systems against manipulation and gaming
- Reviewer assignment, calibration, and incentives; science-of-science analyses of AI's impact on research

Evaluated on soundness, novelty, and relevance to the workshop scope. Rigorous negative results and careful replications are welcome.

need not be technical, and we explicitly welcome perspectives from philosophy, science studies, law, and policy. Titles should begin with "Position:". Topics include, but are not limited to:

- What counts as a publication when AI scientists produce papers at scale?
- Redesigning peer review, publication criteria, and incentive structures for the AI era
- Credit, authorship, and career incentives for human and hybrid human-AI research
- Governance, funding allocation, and institutional responses to AI-accelerated science
- Ethics and epistemology of AI in the scientific process, including policies for AI use in reviewing and authoring

Evaluated primarily on the clarity of the position, the quality of argument and evidence, significance, and serious engagement with alternative views (not strawmen), rather than on empirical novelty.

Submission Guidelines

Submissions are made through the workshop's [OpenReview page](#). If you are new to OpenReview, please create your account well before the deadline, as moderation of new accounts can take up to two weeks.

Submissions are either **4 pages** (concise technical contributions, extended abstracts) or **8 pages** (more substantial contributions, for example position papers), excluding references, with unlimited appendix; reviewers are only obliged to read the main text. Please use the [NeurIPS 2026 LaTeX style file](#), with the footnote changed to "Submitted to AI for Meta-Science workshop (NeurIPS 2026)". The NeurIPS paper checklist is not needed. Submit a single PDF with any appendices included.

Novel work is reviewed double-blind: anonymize the submission, including links and supplementary material, and cite your own prior work in the third person. Previously published work is reviewed single-blind and need not be anonymized. In the spirit of the workshop, we encourage authors to briefly disclose the role AI systems played in producing the submission (for example writing assistance, analysis, or autonomous generation); this is not a requirement, but a short statement is appreciated.

Review Process

Each submission receives at least two reviews (three preferred) on OpenReview. Reviewers must read the full main text and vouch for having done so; they may use AI assistance, for example to check references and appendices, but carry full responsibility for their review. Reviewers grade each paper as reject, accept (poster), or accept (oral), with detailed justification required for reject recommendations; final decisions are made by the organizers. There is no formal rebuttal phase; time permitting, authors of rejected papers may be offered a short response opportunity. The *position track* review pool includes interdisciplinary reviewers from meta-science, philosophy, and policy. Organizers and program committee members adhere to the NeurIPS conflict-of-interest policy.

Archival Status and Dual Submission

The workshop is **non-archival**: accepted papers will be listed on the workshop website and OpenReview, but there are no formal proceedings, and posting to arXiv before or after the workshop is allowed and encouraged. Dual submission is welcome: work under review elsewhere, whether at the NeurIPS 2026 main conference, a journal, or another workshop,

may be submitted here as long as this complies with the other venue’s policy. **Previously published relevant work is also welcome.**

Presentation and Attendance

All accepted papers are presented as posters, and a subset will be selected for short contributed talks. At least one author is expected to attend the workshop in person and present the work, which requires NeurIPS 2026 workshop registration (separate from main-conference registration). Camera-ready versions, de-anonymized and with an optional extra page, are due in November 2026; the exact date will be announced with notifications.

Important Dates (Anywhere on Earth)

Submissions open	July 27, 2026, on the workshop’s OpenReview page	
Paper submission deadline	August 29, 2026	1 weeks 6 days 11:56:31
Accept/reject notification	September 29, 2026	6 weeks 2 days 11:56:31
Camera-ready deadline	November 2026 (exact date TBA)	
Workshop	December 2026 (exact day TBA) at NeurIPS 2026, Paris	

Frequently Asked Questions

Is a position paper without experiments really in scope?

Yes, that is exactly what the position track is for. Position papers are judged on the strength and clarity of their argument, not on empirical results, and are reviewed with interdisciplinary criteria.

Can non-ML researchers (philosophy, science studies, policy, publishing) submit?

Absolutely. The workshop is sociotechnical by design, and the position track especially welcomes contributions from outside machine learning.

Do I need to attend in person?

At least one author is expected to attend in person and present. If exceptional circumstances prevent this, please [contact us](#).

Can AI systems co-author submissions?

AI assistance is welcome; this is, after all, a workshop about AI in science. We encourage a brief disclosure of the role AI played in producing the submission. Responsibility for the content rests with the human authors.

Call for Reviewers

We are recruiting reviewers for both tracks; each reviewer will receive at most 3 papers. If you would like to help, please fill in the [reviewer sign-up form](#).