

WORKSHOP ON

Evaluation of Interactive Agents

NeurIPS 2026

December 12 or 13, 2026
Atlanta, Georgia

Submissions are open now! Submit on OpenReview [here](#). Deadline: August 29, 2026 (AOE).

Under construction: finalized date and call for papers coming soon!

Why this workshop?

The rapid transition from large language models (LLMs) as single-turn assistants to interactive agents has created an urgent need for new evaluation methodologies. LLMs are increasingly deployed in high-impact settings such as education, counseling, negotiation, research assistance, and software development, where success depends not only on generating a correct response, but on sustaining effective interactions over [Skip to content](#) tories.

Evaluating interactive agents directly with real users can be slow, expensive, difficult to reproduce, and hard to scale in expert domains. This has led to growing use of automatic evaluation methods, ranging from rubric-based grading to user simulators, where an LLM simulates user behavior to support evaluation, training, and stress testing. However, many issues with these approaches remain: for example, user simulators may fail to preserve latent user states, reflect diverse human attributes, represent realistic goals, or match the interaction style of real users.

In light of these challenges, this workshop will focus on methods for developing more rigorous, scalable, and scientific evaluation methods for interactive agents.

Topics

We invite contributions on topics including, but not limited to:

- Evaluation protocols for multi-turn assistants, tool-using agents, computer-use agents, collaborative agents, and user-facing systems
 - Trajectory-level evaluation, including transcripts, tool calls, intermediate states, final task outcomes, latency, cost, and other operational metrics
 - Realistic simulation of users, environments, and interaction partners
 - Validation of user simulators as proxies for human behavior and as stress tests for deployed agents
 - Grader design, including deterministic checks, model-based rubrics, human evaluation, and calibration between them
 - Benchmarks for long-horizon interaction, memory, adaptation, error recovery, and reliability across repeated trials
 - Safety, fairness, privacy, and ethical considerations in evaluating
- [Skip to content](#) g from interaction data, production failures, user feedback,

Call for papers

We invite submissions on the topics listed above. Early-stage work is welcome.

- **Where to submit.** Submissions are open now on our [OpenReview submission site](#). The deadline is **August 29, 2026, Anywhere on Earth**.
- **Format and length.** Submissions must use the official NeurIPS 2026 style, available as an [Overleaf template](#). Full papers may be up to 9 pages and short papers up to 4 pages, excluding references and appendices.
- **Double-blind review.** All submissions are reviewed double-blind, so please anonymize your paper: remove author names and affiliations, and avoid identifying information in the text, acknowledgments, and links.
- **Non-archival.** The workshop is non-archival and accepted papers will not appear in published proceedings. You may submit work that is currently under review at, or has already been accepted to, another venue, and you remain free to publish it elsewhere afterwards.
- **Papers accepted to NeurIPS 2026.** Papers already accepted to the NeurIPS 2026 main conference will undergo an expedited review that primarily evaluates their relevance to the workshop themes.
- **Opinion papers.** We welcome opinion papers. The title must state the opinion and follow the format “Opinion: [Your Title]” — for example, “Opinion: Large Language Models Should Not Replace Peer Review in Scientific Publishing”.
- **Presentation.** Accepted work will primarily be presented as posters, with a select number of papers receiving spotlight talks, as well as a best paper award.

[Skip to content](#)

- **In-person attendance.** This is an in-person workshop, and we expect at least one author of each accepted paper to attend and present in person, barring unexpected circumstances.

Important dates	
Submission deadline	August 29, 2026 Anywhere on Earth
Decision notification	September 29, 2026 Anywhere on Earth
Camera-ready deadline	November 6, 2026 Anywhere on Earth
Workshop	December 12 or 13 (TBD), 2026 Atlanta, Georgia


Invited speakers

To Be Confirmed

Schedule

Time	Session
8:00–8:10	Opening remarks
8:10–9:10	Invited talks
9:10–10:00	Oral session
10:00–11:00	Poster and demo session
11:00–12:00	Invited talks
12:00–1:00	Lunch and mentoring tables
1:00–2:00	Invited talks
2:00–3:00	Poster and demo session
3:00–3:15	Break
3:15–4:15	Invited talks
Skip to content	Panel

Organizers



Yao Dou

Georgia Tech



Siyan Li

Columbia University




Marwa Abdulhai

Princeton University



Nicholas Tomlin

TTIC



Michel Galley

Microsoft Research



Jacob Eisenstein

Google DeepMind

[Skip to content](#)



Alan Ritter

Georgia Tech



Wei Xu

Georgia Tech

Questions? Contact the
workshop chairs:

douy@gatech.edu · siyan.li@columbia.edu ·
marwa_abdulhai@berkeley.edu

Workshop on Evaluation of Interactive Agents
NeurIPS 2026 · Atlanta, Georgia

